

Follow the instructions below carefully:

- Only load the **MASS**-package (and no other packages!). From the **MASS**-package you only need the **ginv**-function (with default settings).
- Load the **assignmentA_grpB.Rdata** file with the assignment and group numbers substituted for **A** and **B**, respectively, into R. This file is available via Brightspace.
- Verify that three objects, named **X**, **Y**, and **betas** have been loaded into R's memory. The first two objects contain the data on the response vector (**Y**) and design matrix (**X**). The third one is the true regression parameter β vector.
- Consider the linear regression model $\mathbf{Y} = \mathbf{X}\beta + \varepsilon$ (without intercept) and $\varepsilon \sim \mathcal{N}(\mathbf{0}_n, \sigma^2 \mathbf{I}_{nn})$, to explain the variation in the response **Y** by a linear combination of the columns of the design matrix **X**.
- Fit the linear regression model by means of the ridge regression estimation with penalty parameter $\lambda = B/10$ (with the group number substituted for **B**).
- Extract, in four decimals using the **round**-function from the above evaluated estimator of β , the **B**-th element (with the group number substituted for **B**). This yields the first part of your answer to the exercise that is to be inserted in the email.
- Evaluate the bias of the ridge regression estimate found above, and decompose this bias into a part due to the penalization and one attributable to the high-dimensionality of the design.
- Extract, in four decimals using the **round**-function from the above evaluated high-dimensional attributable part of the bias, the **B**-th element (with the group number substituted for **B**). This yields the second part of your answer to the exercise that is to be inserted in the email.
- Send your answer (accompanied by R-code) in before 23:59 CET, September 23, 2024. Instructions for composing the email can be found in the pdf-file with information on the bonus exercises.