

Statistical Learning Weekly Assignment 05

Joshua Damm


2024-03-14

```
train <- readRDS("masq_train.Rda")
test  <- readRDS("masq_test.Rda")
```

a) Inspect multicollinearity between the numeric MASQ items. What do you expect about relative performance of lasso, ridge and elastic net regression?

```
# select MASQ columns
MASQ = train[, 12:100]

# inspect multicollinearity
hist(cor(MASQ))
```



The pairwise correlations / the correlation matrix follow a normal distribution with a mean around 0.3. There is a slight peak at 1, because of the diagonal entries of the matrix. All 3 techniques might be a feasible solution. Lasso might be an option, since it could remove redundant variables. Elastic net might be the best choice as it combines the benefits of lasso and ridge.

(b) Pick three candidate procedures from ridge, elastic net, lasso, relaxed lasso.

I will choose lasso, ridge and elastic net with an alpha of 0.5.

(c) Use library glmnet to fit the models, and select the most accurate model through 10-fold cross-validation on the training set.

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
x = model.matrix(D_DEPDYS ~ ., data = train)
x_test = model.matrix(D_DEPDYS ~ ., data = test)
y = train$D_DEPDYS
```

```
# setup
```

```
lasso = cv.glmnet(x, y, alpha = 1); lasso
```

```
##
```

```
## Call: cv.glmnet(x = x, y = y, alpha = 1)
```

```
##
```

```
## Measure: Mean-Squared Error
```

```
##
```

| ## | Lambda | Index | Measure | SE | Nonzero |
|--------|---------|-------|---------|----------|---------|
| ## min | 0.01140 | 34 | 0.1721 | 0.005380 | 37 |
| ## 1se | 0.04195 | 20 | 0.1773 | 0.004925 | 17 |

```
ridge = cv.glmnet(x, y, alpha = 0); ridge
```

```
##
```

```
## Call: cv.glmnet(x = x, y = y, alpha = 0)
```

```
##
```

```
## Measure: Mean-Squared Error
```

```
##
```

| ## | Lambda | Index | Measure | SE | Nonzero |
|--------|--------|-------|---------|----------|---------|
| ## min | 0.276 | 74 | 0.1721 | 0.004654 | 132 |
| ## 1se | 1.774 | 54 | 0.1763 | 0.003817 | 132 |

```
elastic = cv.glmnet(x, y, alpha = 0.5); elastic
```

```
##
## Call: cv.glmnet(x = x, y = y, alpha = 0.5)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.01725    37  0.1712 0.005562      49
## 1se 0.08389    20  0.1767 0.004946      19

# predict models
predict_L = predict(lasso, s = "lambda.min", newx = x_test)
predict_R = predict(ridge, s = "lambda.min", newx = x_test)
predict_E = predict(elastic, s = "lambda.min", newx = x_test)

# get mean squared errors
MSE_L = mean((predict_L - test$D_DEPDYS)^2)
MSE_R = mean((predict_R - test$D_DEPDYS)^2)
MSE_E = mean((predict_E - test$D_DEPDYS)^2)

# benchmark
MSE_MAX = var(test$D_DEPDYS)

# cross-validated R2
CVR2_L = 1 - MSE_L/MSE_MAX; CVR2_L
```

```
## [1] 0.3331148
```

```
CVR2_R = 1 - MSE_R/MSE_MAX; CVR2_R
```

```
## [1] 0.3308643
```

```
CVR2_E = 1 - MSE_E/MSE_MAX; CVR2_E
```

```
## [1] 0.334857
```

Lasso shows the best performance with a cross-validated R2 of 0.3337474.

(d) Compute the misclassification rate (MCR) on the test set.

```
preds_L_1se = predict(lasso, newx = x[x_test, ], type = "response")
preds_L_min = predict(lasso, newx = x[x_test, ], type = "response",
                      s = "lambda.min")

tab_L_1se = prop.table(table(preds_L_1se > .5, y[x_test]))
tab_L_min = prop.table(table(preds_L_min > .5, y[x_test]))
tab_L_1se; tab_L_min

##
##           0           1
## FALSE 0.141621916 0.001009058
## TRUE  0.001994920 0.855374105
```

```
##
##           0           1
## FALSE 0.1416857073 0.0007538942
## TRUE  0.0019311289 0.8556292697
```

```
sum(diag(tab_L_1se)); sum(diag(tab_L_min))
```

```
## [1] 0.996996
```

```
## [1] 0.997315
```

For the `sel_lambda` criteria the MCR is a little higher.

(e) Use the `coef` method to extract the selected variables and their coefficients from the best-performing model.

```
L_coefs = coef(lasso, s = "lambda.min")
L_coefs[L_coefs[,1] != 0,]
```

```
## (Intercept)      GENDerv      Leeftijd      DEMOG26      DEMOG32
## -4.814769e-01  8.170714e-03  1.149901e-03 -7.256811e-02  1.328822e-02
##      DEMOG34      DEMOG3NA      DEMOG53      DEMOG55      DEMOG62
## -6.421318e-05 -3.518655e-02 -3.346312e-02  1.993973e-02  5.647645e-02
##      MASQ01      MASQ02      MASQ03      MASQ05      MASQ13
##  3.481755e-02 -1.548544e-02 -6.228674e-03  7.338993e-04  8.940010e-03
##      MASQ14      MASQ16      MASQ18      MASQ21      MASQ22
##  2.865788e-03  6.948313e-02  3.406095e-03  3.920310e-03  2.403443e-02
##      MASQ24      MASQ29      MASQ30      MASQ31      MASQ33
##  5.460323e-03  7.899327e-04  2.265849e-02  8.409185e-03  2.287192e-03
##      MASQ37      MASQ38      MASQ41      MASQ43      MASQ54
##  1.893545e-02  5.152496e-03  2.533319e-02  5.592111e-03  1.655632e-03
##      MASQ59      MASQ60      MASQ62      MASQ70      MASQ76
## -1.041407e-02  9.627503e-03  1.395614e-02  4.299627e-03  9.622140e-03
##      MASQ78      MASQ89      MASQ90
##  9.813473e-03  2.667092e-02  1.438215e-02
```

```
Anhedonic_Depression = c(1, 14, 18, 21, 23, 26, 27, 30, 33, 35, 36, 39, 40, 44, 49, 53, 58, 66, 72, 78,
Anxious_Arousal = c(3, 19, 25, 45, 48, 52, 55, 57, 61, 67, 69, 73, 75, 79, 85, 87, 88)
General_Distress_Depression = c(6, 8, 10, 13, 16, 22, 24, 42, 47, 56, 64, 74)
General_Distress_Anxiety = c(2, 9, 12, 15, 20, 59, 63, 65, 77, 81, 82)
General_Distress_Mixed = c(4, 5, 17, 29, 31, 34, 37, 50, 51, 70, 76, 80, 83, 84, 90)
```